

Assignment: Final Project

Spring 2025 Advanced Machine Learning (BA-64061-001)

Student Name: Sanjar Turaev

Student ID: 811360569

Student e-mail: sturaev@kent.edu

Deep Learning Applications: A Personal Exploration into Computer Vision

1. Application Choice: Computer Vision

When it came time to pick a specific application of deep learning to study, I didn't have to think too hard. For me, it was always going to be computer vision. There's something incredibly fascinating about the idea of teaching a machine to "see" the world and interpret it the way we humans do. Visual information is so fundamental to our everyday lives, we rely on our eyes more than any other sense and helping computers develop this ability feels like a major step forward for technology and society.

Computer vision is everywhere now. It's used in facial recognition, self-driving cars, medical diagnostics, security systems, augmented reality, and even social media apps that apply funny filters. It's so widespread that sometimes we don't even realize how much we interact with it. Personally, I've always been curious about how my phone unlocks with my face or how my favorite apps can identify objects in photos. That curiosity led me straight to computer vision for this project.

At the heart of computer vision's progress is deep learning, particularly Convolutional Neural Networks (CNNs). These specialized models revolutionized the field by learning patterns automatically, instead of relying on humans to define features. However, even as I read more about CNNs, I discovered an exciting new development that is Vision Transformers (ViTs). Unlike CNNs, ViTs treat an image like a series of patches and analyze it more like reading a sequence, much like how language models handle text. This new method has opened up fresh possibilities for better accuracy and flexibility, especially when dealing with huge datasets.

One thing I quickly realized, though, is that teaching machines to "see" isn't simple. Lighting conditions, object occlusions, different angles, all of these challenges can trip up a machine. Even a slight alteration, like a sticker on a stop sign, can confuse an AI-powered self-driving car. That

made me realize just how nuanced and difficult true computer vision really is. Machines don't naturally interpret images the way humans do; they need massive datasets and careful training.

Still, the potential is undeniable. In healthcare, AI models are saving lives by identifying tumors early. In transportation, they're making cars safer and smarter. In security, they're helping spot threats before they escalate. It's clear that computer vision is no longer just an academic exercise, it's a critical tool for real-world applications.

Choosing computer vision wasn't just about finding a cool topic. It was about exploring a field that bridges human perception and machine intelligence in ways that could have profound impacts on our future. As someone who is passionate about both technology and its real-world applications, computer vision felt like the perfect fit for this deep learning project. I wanted to understand not just how machines learn to see, but how that vision is reshaping industries, solving real problems, and creating new ones. And honestly, the more I researched, the more I realized just how much potential and how many challenges this field still holds.

2. Literature Review: What's Happening in Computer Vision Today

Over the past decade, deep learning has dramatically changed the field of computer vision. Early computer vision models used manually designed features like edges or corners. However, the introduction of Convolutional Neural Networks (CNNs) marked a major shift. CNNs automatically learn important features directly from images, reducing the need for human intervention and achieving much better accuracy.

One of the biggest milestones was the 2012 ImageNet competition, where the CNN model AlexNet, created by Krizhevsky et al., beat all competitors by a large margin. Since then, CNNs like VGGNet, ResNet, and DenseNet have continued to improve image classification, object detection, and segmentation tasks. ResNet, in particular, introduced the concept of residual connections, which made it possible to train very deep networks without suffering from the "vanishing gradient" problem.

More recently, Vision Transformers (ViT) have emerged as an alternative to CNNs. Inspired by the success of transformers in natural language processing (like BERT and GPT), ViTs treat an image as a series of patches, similar to how a sentence is a sequence of words. This allows the model to capture relationships between different parts of an image without relying on convolutional filters. Research has shown that ViTs perform better than CNNs when trained on very large datasets, although they tend to require more data and computational resources.

Another important area of research is self-supervised learning. Instead of needing millions of labeled images, self-supervised methods allow models to learn useful representations from unlabeled data by creating tasks where the data itself provides the supervision. For example, a model might learn to predict the rotation of an image or fill in missing parts. Self-supervised

approaches like SimCLR, MoCo, and DINO have shown impressive results and are likely to play a big role in the future of computer vision.

However, despite these advances, deep learning models still have some major weaknesses. They can be fooled by adversarial examples, small changes to an image that cause the model to make wrong predictions. They also often struggle to generalize to new situations they haven't seen before. Moreover, many models are "black boxes," meaning it is hard to understand why they make certain decisions.

Researchers are now working on ways to make models more robust, interpretable, and efficient. Techniques like attention mechanisms, model distillation, and neural architecture search are helping create models that are smaller, faster, and more understandable.

Overall, the field of computer vision is very dynamic. Each year brings new methods and better performance, but also reveals new challenges. Keeping up with this fast-moving field requires a strong understanding of the core ideas, as well as an openness to new techniques and tools that may replace today's best models.

3. Industry Applications: How Computer Vision is Changing the World

Deep learning-powered computer vision is being adopted in many industries, transforming how we work, live, and interact with technology. Let's look at how different sectors are using it today and where it might go next.

Healthcare is one of the biggest beneficiaries. Deep learning models are used to analyze medical images like X-rays, CT scans, and MRIs. For instance, companies like Aidoc and Zebra Medical Vision have developed AI tools that can spot diseases such as cancer, pneumonia, or brain hemorrhages faster and sometimes more accurately than human doctors. AI can also assist doctors in prioritizing urgent cases, saving valuable time in critical situations. Another exciting area is digital pathology, where AI systems can help examine tissue samples for abnormalities.

In **transportation**, computer vision plays a key role in autonomous vehicles. Self-driving cars from companies like Tesla, Waymo, and Cruise use deep learning models to understand the road environment, identifying other vehicles, pedestrians, traffic signs, and lane markings. Computer vision also supports advanced driver-assistance systems (ADAS) like automatic emergency braking and lane-keeping assistance, which are already available in many modern cars.

Security and surveillance are also major fields where computer vision is growing. Facial recognition technology, although controversial, is widely used in airports, border security, and law enforcement to identify individuals quickly. AI-based video analytics can detect suspicious behavior, count people, monitor social distancing (important during pandemics), and even spot abandoned objects in public places.

In **retail**, AI vision systems are used for inventory management, cashier-less stores (like Amazon Go), customer tracking, and personalized marketing. Smart cameras can automatically restock shelves, monitor theft, and optimize store layouts based on customer movement patterns.

Manufacturing industries use computer vision for quality control. AI models inspect products for defects faster and more consistently than human workers, reducing waste and improving production efficiency. Drones equipped with computer vision are being used for infrastructure inspection, agriculture (monitoring crop health), and disaster response.

While these examples show the massive potential of computer vision, there are also concerns. Issues around data privacy, bias in facial recognition, and the reliability of AI in safety critical applications must be carefully addressed. Additionally, the high cost of developing and deploying deep learning models means that small businesses may struggle to adopt these technologies without external support.

Despite these challenges, the future looks bright. As models become cheaper to train and easier to interpret, we can expect computer vision to become a standard tool in more industries, helping to automate tasks, improve accuracy, and open new possibilities that were once considered science fiction.

4. Future Developments: What's Next for Computer Vision

Even though computer vision has made incredible progress, it's clear from my research that the field still has many challenges and exciting opportunities ahead. Personally, I think we're standing at the edge of a new era for visual AI.

One of the biggest problems right now is that deep learning models need huge datasets to perform well. Gathering and labeling millions of images is expensive, time-consuming, and often impractical, especially for smaller companies. That's why I'm particularly excited about self supervised and few-shot learning approaches. If we can teach models to learn more like humans making sense of a small number of examples, we'll unlock a new level of accessibility and creativity. I imagine a future where startups and small labs can compete with tech giants without needing giant data warehouses.

Another major hurdle is robustness. Right now, it's disturbingly easy to trick an AI model with slight changes to an image. An autonomous car mistaking a stop sign for a speed limit sign because of a few stickers is not just a funny story, it could be deadly. Developing models that can reason about what they see, not just memorize patterns, feels essential for making computer vision safe and trustworthy.

Explainability is another issue close to my heart. I think a lot of the public distrust of AI comes from the fact that most deep learning models are black boxes. Even engineers often don't fully understand why a model made a certain decision. Research into explainable AI (XAI) is starting to help, using tools like Grad-CAM, LIME, and SHAP, but there's still a long way to go. In the future, I hope that users whether doctors, drivers, or store managers will be able to understand and challenge AI decisions easily.

Energy consumption is another hidden cost. Training massive models like Vision Transformers can require weeks of GPU time and tons of electricity. It's not sustainable in the long run. I'm hopeful that more efficient models, like EfficientNet or techniques like model pruning and quantization, will make computer vision greener. Maybe someday we'll even have state-of-the-art models that run on mobile devices without needing a cloud server.

Finally, I think combining computer vision with other senses and types of reasoning will be the next big leap. Instead of treating images alone, AI could integrate visual data with sound, touch, and even symbolic reasoning to understand the world more holistically kind of like how humans use multiple senses to make better decisions.

Looking ahead, I feel both excited and cautious. We have powerful tools in our hands, but how we choose to use them and how carefully we design them will determine whether computer vision becomes a force for good or a source of new problems. Either way, it's clear that this field will continue to shape our future in profound ways.

Conclusion

Working on this project has truly changed how I think about computer vision. Before starting, I saw it mainly as a cool technology for futuristic gadgets. Now, I realize it's something much deeper and more important. Computer vision is transforming industries, saving lives, and raising big ethical questions, all at the same time.

From the breakthroughs of CNNs to the exciting possibilities of Vision Transformers and self-supervised learning, the field is advancing faster than I ever imagined. Yet it's not without its challenges. Issues like model fragility, the need for massive datasets, lack of transparency, and environmental impact all show that there's still a long road ahead. It's humbling to realize how complex "seeing" really is something we humans take for granted every day.

What excites me most is how many opportunities there are to improve things. Whether it's making models more robust, more explainable, or more energy efficient, there's so much important work to be done. I feel motivated to be part of that future helping create AI systems that are not just powerful, but also safe, fair, and accessible to everyone.

This project has reinforced for me that computer vision isn't just about technology. It's about bridging the gap between machines and the human world, and doing it in a way that benefits society. I'm grateful for the chance to dive deep into this topic, and even more excited to keep learning and contributing to this amazing field in the future.

References

- Krizhevsky, A., Sutskever, I., & Hinton, G. E. (2012). "ImageNet Classification with Deep Convolutional Neural Networks." *Advances in Neural Information Processing Systems*.

- He, K., Zhang, X., Ren, S., & Sun, J. (2016). "Deep Residual Learning for Image Recognition." *Proceedings of the IEEE Conference on Computer Vision and Pattern Recognition (CVPR)*.
- Dosovitskiy, A., et al. (2020). "An Image is Worth 16x16 Words: Transformers for Image Recognition at Scale." *International Conference on Learning Representations (ICLR)*.
- Chen, T., Kornblith, S., Norouzi, M., & Hinton, G. (2020). "A Simple Framework for Contrastive Learning of Visual Representations." *International Conference on Machine Learning (ICML)*.
- Redmon, J., & Farhadi, A. (2018). "YOLOv3: An Incremental Improvement." *arXiv preprint*.